

```
In [1]: import sklearn
import numpy as np
```

```
In [41]: from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras.wrappers.scikit_learn import KerasClassifier
```

```
In [63]: from sklearn.linear_model import LogisticRegression
```

```
In [2]: from keras.datasets import imdb
(training_data, training_targets), (testing_data, testing_targets) = imdb.load_data
```

Using TensorFlow backend.

```
c:\program files\python37\lib\site-packages\tensorflow\python\framework\dtypes.
py:526: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is dep
recated; in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)])
```

```
c:\program files\python37\lib\site-packages\tensorflow\python\framework\dtypes.
py:527: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is dep
recated; in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)])
```

```
c:\program files\python37\lib\site-packages\tensorflow\python\framework\dtypes.
py:528: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is dep
recated; in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)])
```

```
c:\program files\python37\lib\site-packages\tensorflow\python\framework\dtypes.
py:529: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is dep
recated; in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)])
```

```
c:\program files\python37\lib\site-packages\tensorflow\python\framework\dtypes.
py:530: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is dep
recated; in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
```

```
_np_qint32 = np.dtype [("qint32", np.int32, 1)])
```

```
c:\program files\python37\lib\site-packages\tensorflow\python\framework\dtypes.
py:535: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is dep
recated; in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
```

```
np_resource = np.dtype [("resource", np.ubyte, 1)])
```

```
In [3]: training_data.shape
```

```
Out[3]: (25000,)
```

```
In [5]: training_targets
```

```
Out[5]: array([1, 0, 0, ..., 0, 1, 0], dtype=int64)
```

```
In [6]: np.count_nonzero(training_targets)
```

```
Out[6]: 12500
```

```
In [7]: print("Categories:", np.unique(training_targets))
```

```
Categories: [0 1]
```

```
In [8]: print("Number of unique words:", len(np.unique(np.hstack(training_data))))
```

```
Number of unique words: 9998
```

```
In [9]: index = imdb.get_word_index()
reverse_index = dict([(value, key) for (key, value) in index.items()])
```

```
In [10]: decoded = " ".join( [reverse_index.get(i - 3, "#") for i in training_data[0]] )
print(decoded)
```

```
# this film was just brilliant casting location scenery story direction everyone's really suited the part they played and you could just imagine being there robert # is an amazing actor and now the same being director # father came from the same scottish island as myself so i loved the fact there was a real connection with this film the witty remarks throughout the film were great it was just brilliant so much that i bought the film as soon as it was released for # and would recommend it to everyone to watch and the fly fishing was amazing really cried at the end it was so sad and you know what they say if you cry at a film it must have been good and this definitely was also # to the two little boy's that played the # of norman and paul they were just brilliant children are often left out of the # list i think because the stars that play them all grown up are such a big profile for the whole film but these children are amazing and should be praised for what they have done don't you think the whole story was so lovely because it was true and was someone's life after all that was shared with us all
```

```
In [11]: from keras.preprocessing import sequence
```

```
In [12]: X_tr = sequence.pad_sequences(training_data, maxlen=100)
#X_ts = sequence.pad_sequences(testing_data, maxlen=100)
```

```
In [179]: X_train_1 = []

for i in range(0, len(X_tr)):

    decoded = " ".join( [reverse_index.get(i - 3, "#") for i in X_tr[i]] )
    X_train_1.append(decoded)
```

```
In [14]: #X_test = []

#for i in range(0,len(X_ts)):

#    decoded = " ".join( [reverse_index.get(i - 3, "#") for i in X_ts[i]] )
#    X_test.append(decoded)
```

```
In [180]: X_train_1
```

```
Out[180]: ["cry at a film it must have been good and this definitely was also # to the
two little boy's that played the # of norman and paul they were just brillian
t children are often left out of the # list i think because the stars that pl
ay them all grown up are such a big profile for the whole film but these chil
dren are amazing and should be praised for what they have done don't you thin
k the whole story was so lovely because it was true and was someone's life af
ter all that was shared with us all",
"funny in equal # the hair is big lots of boobs # men wear those cut # shirt
s that show off their # sickening that men actually wore them and the music i
s just # trash that plays over and over again in almost every scene there is
trashy music boobs and # taking away bodies and the gym still doesn't close f
or # all joking aside this is a truly bad film whose only charm is to look ba
ck on the disaster that was the 80's and have a good old laugh at how bad eve
rything was back then",
"touching the floor at how bad it really was the rest of the time everyone e
lse in the theatre just started talking to each other leaving or generally cr
ying into their popcorn that they actually paid money they had # working to w
atch this feeble excuse for a film it must have looked like a great idea on p
aper but on film it looks like no one in the film has a clue what is going on
"
```

CountVectorizer and Tfidf

```
In [16]: from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
import nltk
```

```
In [17]: vec1 = CountVectorizer(min_df=2, tokenizer=nltk.word_tokenize)
```

```
In [18]: train_cv = vec1.fit_transform(X_train_1)
```

```
In [19]: train_cv
```

```
Out[19]: <25000x9815 sparse matrix of type '<class 'numpy.int64'>'
with 1706648 stored elements in Compressed Sparse Row format>
```

```
In [20]: vec1.vocabulary_.get('good')
```

```
Out[20]: 3889
```

```
In [21]: vec1.vocabulary_.get('hot')
```

```
Out[21]: 4329
```

```
In [22]: train_cv.shape
```

```
Out[22]: (25000, 9815)
```

```
In [61]: train_cv.shape
```

```
Out[61]: (25000, 9815)
```

```
In [64]: X_train_cv = train_cv[:20000]
y_train_cv = training_targets[:20000]
X_test_cv = train_cv[20000::]
y_test_cv = training_targets[20000::]
```

```
In [65]: print("X_train_cv shape: ", train_cv[:20000].shape)
print("y_train_cv shape: ", training_targets[:20000].shape)
print("X_test_cv shape: ", train_cv[20000::].shape)
print("y_test_cv shape: ", training_targets[20000::].shape)
```

```
X_train_cv shape: (20000, 9815)
y_train_cv shape: (20000,)
X_test_cv shape: (5000, 9815)
y_test_cv shape: (5000,)
```

Logistic Regression with CountVectorizer

```
In [66]: classifier = LogisticRegression()
classifier.fit(X_train_cv, y_train_cv)
```

```
c:\program files\python37\lib\site-packages\sklearn\linear_model\logistic.py:43
2: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a
solver to silence this warning.
FutureWarning)
```

```
Out[66]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='warn', n_jobs=None, penalty='l2',
random_state=None, solver='warn', tol=0.0001, verbose=0,
warm_start=False)
```

```
In [68]: print("Logistic Regression on Count Vectorizer Accuracy = ", classifier.score(X_t
```

```
Logistic Regression on Count Vectorizer Accuracy = 82.74000000000001
```

Neural Networks with Count Vectorizer

```
In [78]: def build_model():  
    model = Sequential()  
    model.add(Dense(512, input_dim=9815, activation='elu'))  
    model.add(Dropout(0.3))  
    model.add(Dense(1, activation='sigmoid'))  
    model.compile(loss='binary_crossentropy', optimizer='sgd', metrics=['accuracy'])  
    model.summary()  
    return model
```

```
In [80]: cv_estimator = KerasClassifier(build_fn=build_model, epochs=25, batch_size=64)
cv_estimator.fit(x = X_train_cv, y = y_train_cv, validation_split=0.25, workers=2)
```

Model: "sequential_16"

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 512)	5025792
dropout_26 (Dropout)	(None, 512)	0
dense_42 (Dense)	(None, 1)	513

Total params: 5,026,305

Trainable params: 5,026,305

Non-trainable params: 0

Train on 15000 samples, validate on 5000 samples

Epoch 1/25

15000/15000 [=====] - 5s 364us/step - loss: 0.6416 - accuracy: 0.6337 - val_loss: 0.5871 - val_accuracy: 0.7052

Epoch 2/25

15000/15000 [=====] - 5s 349us/step - loss: 0.5688 - accuracy: 0.7151 - val_loss: 0.5457 - val_accuracy: 0.7376

Epoch 3/25

15000/15000 [=====] - 5s 349us/step - loss: 0.5303 - accuracy: 0.7475 - val_loss: 0.5150 - val_accuracy: 0.7446

Epoch 4/25

15000/15000 [=====] - 5s 349us/step - loss: 0.5028 - accuracy: 0.7634 - val_loss: 0.4907 - val_accuracy: 0.7666

Epoch 5/25

15000/15000 [=====] - 5s 350us/step - loss: 0.4759 - accuracy: 0.7861 - val_loss: 0.4778 - val_accuracy: 0.7712

Epoch 6/25

15000/15000 [=====] - 5s 349us/step - loss: 0.4626 - accuracy: 0.7887 - val_loss: 0.4726 - val_accuracy: 0.7708

Epoch 7/25

15000/15000 [=====] - 5s 349us/step - loss: 0.4503 - accuracy: 0.7968 - val_loss: 0.4641 - val_accuracy: 0.7712

Epoch 8/25

15000/15000 [=====] - 5s 348us/step - loss: 0.4339 - accuracy: 0.8071 - val_loss: 0.4755 - val_accuracy: 0.7632

Epoch 9/25

15000/15000 [=====] - 5s 350us/step - loss: 0.4207 - accuracy: 0.8126 - val_loss: 0.4364 - val_accuracy: 0.7912

Epoch 10/25

15000/15000 [=====] - 5s 351us/step - loss: 0.4128 - accuracy: 0.8141 - val_loss: 0.4387 - val_accuracy: 0.7896

Epoch 11/25

15000/15000 [=====] - 5s 350us/step - loss: 0.4048 - accuracy: 0.8201 - val_loss: 0.4125 - val_accuracy: 0.8038

Epoch 12/25

15000/15000 [=====] - 5s 348us/step - loss: 0.3991 - accuracy: 0.8229 - val_loss: 0.4038 - val_accuracy: 0.8104

Epoch 13/25

15000/15000 [=====] - 5s 349us/step - loss: 0.3896 - accuracy: 0.8301 - val_loss: 0.4036 - val_accuracy: 0.8104

```

Epoch 14/25
15000/15000 [=====] - 5s 349us/step - loss: 0.3793 - a
ccuracy: 0.8366 - val_loss: 0.3941 - val_accuracy: 0.8144
Epoch 15/25
15000/15000 [=====] - 5s 352us/step - loss: 0.3752 - a
ccuracy: 0.8344 - val_loss: 0.3863 - val_accuracy: 0.8222
Epoch 16/25
15000/15000 [=====] - 5s 349us/step - loss: 0.3697 - a
ccuracy: 0.8412 - val_loss: 0.5632 - val_accuracy: 0.7426
Epoch 17/25
15000/15000 [=====] - 5s 349us/step - loss: 0.3605 - a
ccuracy: 0.8455 - val_loss: 0.3858 - val_accuracy: 0.8198
Epoch 18/25
15000/15000 [=====] - 5s 348us/step - loss: 0.3672 - a
ccuracy: 0.8383 - val_loss: 0.4109 - val_accuracy: 0.8114
Epoch 19/25
15000/15000 [=====] - 5s 349us/step - loss: 0.3556 - a
ccuracy: 0.8475 - val_loss: 0.3837 - val_accuracy: 0.8230
Epoch 20/25
15000/15000 [=====] - 5s 350us/step - loss: 0.3428 - a
ccuracy: 0.8564 - val_loss: 0.4431 - val_accuracy: 0.7838
Epoch 21/25
15000/15000 [=====] - 5s 349us/step - loss: 0.3446 - a
ccuracy: 0.8517 - val_loss: 0.3664 - val_accuracy: 0.8336
Epoch 22/25
15000/15000 [=====] - 5s 350us/step - loss: 0.3420 - a
ccuracy: 0.8538 - val_loss: 0.3983 - val_accuracy: 0.8138
Epoch 23/25
15000/15000 [=====] - 5s 350us/step - loss: 0.3388 - a
ccuracy: 0.8586 - val_loss: 0.3746 - val_accuracy: 0.8244
Epoch 24/25
15000/15000 [=====] - 5s 350us/step - loss: 0.3360 - a
ccuracy: 0.8580 - val_loss: 0.3766 - val_accuracy: 0.8224
Epoch 25/25
15000/15000 [=====] - 5s 350us/step - loss: 0.3292 - a
ccuracy: 0.8610 - val_loss: 0.3915 - val_accuracy: 0.8146

```

Out[80]: <keras.callbacks.callbacks.History at 0x169eb55d608>

Comments

As we can see, a logistic regression model performed slightly better than my neural network. I think the reason is my neural network model is overfitting and the accuracy should be very similar to the logistic regression model. I tried a few different neural networks with different attributes but they all performed similarly and there was no improvements compared to the logistic regression.

```
In [ ]: #vec2 = CountVectorizer(min_df=2, tokenizer=nltk.word_tokenize)
```

```
In [ ]: #test_cv = vec2.fit_transform(X_test)
```

```
In [ ]: #test_cv
```

```
In [ ]: #vec2.vocabulary_.get('good')
```

```
In [ ]: #vec2.vocabulary_.get('hot')
```

```
In [ ]: #test_cv.shape
```

Tfidf

- <https://towardsdatascience.com/machine-learning-nlp-text-classification-using-scikit-learn-python-and-nltk-c52b92a7c73a> (<https://towardsdatascience.com/machine-learning-nlp-text-classification-using-scikit-learn-python-and-nltk-c52b92a7c73a>)

```
In [24]: tfidf_transformer = TfidfTransformer()
```

```
In [25]: train_tfidf = tfidf_transformer.fit_transform(train_cv)
#test_tfidf = tfidf_transformer.fit_transform(test_cv)
```

```
In [26]: train_tfidf
```

```
Out[26]: <25000x9815 sparse matrix of type '<class 'numpy.float64'>'
         with 1706648 stored elements in Compressed Sparse Row format>
```

```
In [ ]: #print(test_tfidf)
```

```
In [27]: train_tfidf.shape
```

```
Out[27]: (25000, 9815)
```

```
In [33]: X_train_tfidf = train_tfidf[:20000]
y_train = training_targets[:20000]
X_test_tfidf = train_tfidf[20000::]
y_test = training_targets[20000::]
```

```
In [34]: print("X_train_tfidf shape: ", train_tfidf[:20000].shape)
print("y_train shape: ", training_targets[:20000].shape)
print("X_test_tfidf shape: ", train_tfidf[20000::].shape)
print("y_test shape: ", training_targets[20000::].shape)
```

```
X_train_tfidf shape: (20000, 9815)
y_train shape: (20000,)
X_test_tfidf shape: (5000, 9815)
y_test shape: (5000,)
```

```
In [ ]: #test_tfidf.shape
```

A Naive Bayes model with Tfidf


```
In [35]: from sklearn.naive_bayes import MultinomialNB  
clf = MultinomialNB().fit(X_train_tfidf, y_train)
```

```
In [39]: predicted = clf.predict(X_test_tfidf)  
print( "Accuracy of Multinomial Nasive Bayes = ",(np.mean(predicted == y_test)*100)  
  
Accuracy of Multinomial Nasive Bayes = 83.0
```

Neural Networks for TF-IDF

```
In [56]: def build_model():  
    model = Sequential()  
    model.add(Dense(512, input_dim=9815, activation='elu'))  
    model.add(Dropout(0.4))  
    model.add(Dense(1, activation='sigmoid'))  
    model.compile(loss='binary_crossentropy', optimizer='adadelta', metrics=['acc'])  
    model.summary()  
    return model
```

```
In [59]: estimator = KerasClassifier(build_fn=build_model, epochs=10, batch_size=128)
estimator.fit(x = X_train_tfidf, y = y_train, validation_split=0.2, workers=2)
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
dense_26 (Dense)	(None, 512)	5025792
dropout_17 (Dropout)	(None, 512)	0
dense_27 (Dense)	(None, 1)	513
Total params: 5,026,305		
Trainable params: 5,026,305		
Non-trainable params: 0		

Train on 16000 samples, validate on 4000 samples

Epoch 1/10

16000/16000 [=====] - 8s 515us/step - loss: 0.6452 - accuracy: 0.6856 - val_loss: 0.5721 - val_accuracy: 0.7605

Epoch 2/10

16000/16000 [=====] - 8s 503us/step - loss: 0.4850 - accuracy: 0.8068 - val_loss: 0.4316 - val_accuracy: 0.8058

Epoch 3/10

16000/16000 [=====] - 8s 504us/step - loss: 0.3801 - accuracy: 0.8461 - val_loss: 0.3867 - val_accuracy: 0.8305

Epoch 4/10

16000/16000 [=====] - 8s 505us/step - loss: 0.3245 - accuracy: 0.8714 - val_loss: 0.3479 - val_accuracy: 0.8415

Epoch 5/10

16000/16000 [=====] - 8s 504us/step - loss: 0.2923 - accuracy: 0.8845 - val_loss: 0.3355 - val_accuracy: 0.8518

Epoch 6/10

16000/16000 [=====] - 8s 503us/step - loss: 0.2634 - accuracy: 0.8995 - val_loss: 0.3353 - val_accuracy: 0.8503

Epoch 7/10

16000/16000 [=====] - 8s 503us/step - loss: 0.2404 - accuracy: 0.9099 - val_loss: 0.3303 - val_accuracy: 0.8587

Epoch 8/10

16000/16000 [=====] - 8s 505us/step - loss: 0.2222 - accuracy: 0.9174 - val_loss: 0.3350 - val_accuracy: 0.8593

Epoch 9/10

16000/16000 [=====] - 8s 504us/step - loss: 0.2061 - accuracy: 0.9247 - val_loss: 0.3425 - val_accuracy: 0.8535

Epoch 10/10

16000/16000 [=====] - 8s 504us/step - loss: 0.1929 - accuracy: 0.9325 - val_loss: 0.3470 - val_accuracy: 0.8558

Out[59]: <keras.callbacks.callbacks.History at 0x169846ea648>

Comments

The results of Tfidf make more sense compared to the CountVectorizer. The Naive Bayes model performed about 83% after applying Tfidf on dataset. It's slightly better than the performance of CountVectorizer with both neural networks or logistic regression. Also, we can see that the neural network performed better than the Naive Bayes model in Tfidf. Even though we can see that the neural network is against overfitting. I tried different neural networks but they all performed as good as 86%.

In []:

GENSIM

```
In [81]: import gensim
from gensim import utils
from gensim.models.doc2vec import LabeledSentence
from gensim.models import Doc2Vec
```

Doc2Vec

```
In [137]: from gensim.test.utils import common_texts
from gensim.models.doc2vec import Doc2Vec, TaggedDocument
from sklearn.metrics import accuracy_score, f1_score
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn import utils
import csv
from tqdm import tqdm
import multiprocessing
import nltk
from nltk.corpus import stopwords
```

```
In [192]: def read_corpus(doc):
            for i, line in enumerate(doc):

                tokens = gensim.utils.simple_preprocess(line)

                yield gensim.models.doc2vec.TaggedDocument(tokens, str(training_target))
```

```
In [193]: corp = list(read_corpus(X_train_1))
```



```
In [225]: gen_estimator = KerasClassifier(build_fn=build_model, epochs=5, batch_size=32)
gen_estimator.fit(x = np.asarray(gen_X_train), y = np.asarray(gen_y_train), validate
```

Model: "sequential_22"

Layer (type)	Output Shape	Param #
dense_65 (Dense)	(None, 1024)	103424
dropout_44 (Dropout)	(None, 1024)	0
dense_66 (Dense)	(None, 1024)	1049600
dropout_45 (Dropout)	(None, 1024)	0
dense_67 (Dense)	(None, 256)	262400
dropout_46 (Dropout)	(None, 256)	0
dense_68 (Dense)	(None, 256)	65792
dropout_47 (Dropout)	(None, 256)	0
dense_69 (Dense)	(None, 128)	32896
dropout_48 (Dropout)	(None, 128)	0
dense_70 (Dense)	(None, 1)	129
Total params: 1,514,241		
Trainable params: 1,514,241		
Non-trainable params: 0		

Train on 15000 samples, validate on 5000 samples

Epoch 1/5

15000/15000 [=====] - 7s 489us/step - loss: 0.4024 - accuracy: 0.8179 - val_loss: 0.3011 - val_accuracy: 0.8786

Epoch 2/5

15000/15000 [=====] - 7s 457us/step - loss: 0.3192 - accuracy: 0.8651 - val_loss: 0.3984 - val_accuracy: 0.8402

Epoch 3/5

15000/15000 [=====] - 7s 457us/step - loss: 0.3112 - accuracy: 0.8729 - val_loss: 0.3242 - val_accuracy: 0.8604

Epoch 4/5

15000/15000 [=====] - 7s 456us/step - loss: 0.3021 - accuracy: 0.8751 - val_loss: 0.3116 - val_accuracy: 0.8606

Epoch 5/5

15000/15000 [=====] - 7s 457us/step - loss: 0.2926 - accuracy: 0.8786 - val_loss: 0.2980 - val_accuracy: 0.8730

Out[225]: <keras.callbacks.callbacks.History at 0x16a70f714c8>

Comments

We can see that the logistic regression with Gensim did not performed very good. Maybe some attributes could be tuned or other machine learning models like Naive Bayes or SVM could be used to improve it. But the neural network model performed better than other vectorizers and models about 87%. I tried other neural networks too but the best performance so far was 87%. Based on my experience from the last two models with Tfidf and CountVectorizer I only trained my model for 5 epochs to prevent overfitting and it seems that the model isnt overfitting by looking at the val_loss.

Conclusion

As a conclusion, Neural network performed better than the machine learning classifiers except for CountVectorizer. Among the three vectorizers, Gensim performed the best with 87% and CountVectorizer had the poorest performance about 81%.

My failed attempts:

```
In [85]: import nltk as nl
```

```
In [189]: def read_corpus(doc, tokens_only=False):
            for i, line in enumerate(doc):
                tokens = gensim.utils.simple_preprocess(line)

                if tokens_only:
                    yield tokens
                else:
                    # For training data, add tags
                    yield gensim.models.doc2vec.TaggedDocument(tokens, str(training_t
```

```
In [89]: np.array(X_train_1).shape
```

```
Out[89]: (25000,)
```

```
In [90]: X_train_gen = X_train_1[:20000]
          y_train_gen = training_targets[:20000]
          X_test_gen = X_train_1[20000::]
          y_test_gen = training_targets[20000::]
```

```
In [109]: print(np.array(X_train_gen).shape)
          print(np.array(X_test_gen).shape)
```

```
(20000,)
(5000,)
```

```
In [190]: train_corpus = list(read_corpus(X_train_gen))
          test_corpus = list(read_corpus(X_test_gen, tokens_only=True))
```

In [191]: train_corpus

Out[191]: [TaggedDocument(words=['cry', 'at', 'film', 'it', 'must', 'have', 'been', 'go
od', 'and', 'this', 'definitely', 'was', 'also', 'to', 'the', 'two', 'littl
e', 'boy', 'that', 'played', 'the', 'of', 'norman', 'and', 'paul', 'they', 'w
ere', 'just', 'brilliant', 'children', 'are', 'often', 'left', 'out', 'of',
'the', 'list', 'think', 'because', 'the', 'stars', 'that', 'play', 'them', 'a
ll', 'grown', 'up', 'are', 'such', 'big', 'profile', 'for', 'the', 'whole',
'film', 'but', 'these', 'children', 'are', 'amazing', 'and', 'should', 'be',
'praised', 'for', 'what', 'they', 'have', 'done', 'don', 'you', 'think', 'th
e', 'whole', 'story', 'was', 'so', 'lovely', 'because', 'it', 'was', 'true',
'and', 'was', 'someone', 'life', 'after', 'all', 'that', 'was', 'shared', 'wi
th', 'us', 'all'], tags='1'),
TaggedDocument(words=['funny', 'in', 'equal', 'the', 'hair', 'is', 'big', 'l
ots', 'of', 'boobs', 'men', 'wear', 'those', 'cut', 'shirts', 'that', 'show',
'off', 'their', 'sickening', 'that', 'men', 'actually', 'wore', 'them', 'an
d', 'the', 'music', 'is', 'just', 'trash', 'that', 'plays', 'over', 'and', 'o
ver', 'again', 'in', 'almost', 'every', 'scene', 'there', 'is', 'trashy', 'mu
sic', 'boobs', 'and', 'taking', 'away', 'bodies', 'and', 'the', 'gym', 'stil
l', 'doesn', 'close', 'for', 'all', 'joking', 'aside', 'this', 'is', 'truly',
'bad', 'film', 'whose', 'only', 'charm', 'is', 'to', 'look', 'back', 'on', 't

In [141]: np.array(train_corpus).shape

Out[141]: (20000, 2)

In [93]: test_corpus

Out[93]: [['was',
'probably',
'creature',
'ms',
'is',
'unfortunately',
'not',
'werewolf',
'she',
'is',
'merely',
'very',
'strong',
'lunatic',
'br',
'br',
'as',
'film',
'werewolf',
'woman']

In [96]: import random

In [94]: model = Doc2Vec()

model.build_vocab(train_corpus)


```
In [97]: for epoch in range(10):
            model.train(
                train_corpus, total_examples= model.corpus_count,
                epochs=model.epochs)
            # shuffle the corpus
            random.shuffle(train_corpus)
            # decrease the Learning rate
            model.alpha -= 0.0002
            # fix the Learning rate, no decay
            model.min_alpha = model.alpha
```

```
In [ ]: #model.train(train_corpus, total_examples=model.corpus_count, epochs=10)
```

```
In [134]: model.wv.vocab.keys()
```

```
Out[134]: dict_keys(['cry', 'at', 'film', 'it', 'must', 'have', 'been', 'good', 'and',  
    'this', 'definitely', 'was', 'also', 'to', 'the', 'two', 'little', 'boy', 'th  
    at', 'played', 'of', 'norman', 'paul', 'they', 'were', 'just', 'brilliant',  
    'children', 'are', 'often', 'left', 'out', 'list', 'think', 'because', 'star  
    s', 'play', 'them', 'all', 'grown', 'up', 'such', 'big', 'profile', 'for', 'w  
    hole', 'but', 'these', 'amazing', 'should', 'be', 'praised', 'what', 'done',  
    'don', 'you', 'story', 'so', 'lovely', 'true', 'someone', 'life', 'after', 's  
    hared', 'with', 'us', 'funny', 'in', 'equal', 'hair', 'is', 'lots', 'boobs',  
    'men', 'wear', 'those', 'cut', 'shirts', 'show', 'off', 'their', 'sickening',  
    'actually', 'wore', 'music', 'trash', 'plays', 'over', 'again', 'almost', 'ev  
    ery', 'scene', 'there', 'trashy', 'taking', 'away', 'bodies', 'gym', 'still',  
    'doesn', 'close', 'joking', 'aside', 'truly', 'bad', 'whose', 'only', 'char  
    m', 'look', 'back', 'on', 'disaster', 'old', 'laugh', 'how', 'everything', 't  
    hen', 'touching', 'floor', 'really', 'rest', 'time', 'everyone', 'else', 'the  
    atre', 'started', 'talking', 'each', 'other', 'leaving', 'or', 'generally',  
    'crying', 'into', 'popcorn', 'paid', 'money', 'had', 'working', 'watch', 'fee  
    ble', 'excuse', 'looked', 'like', 'great', 'idea', 'paper', 'looks', 'no', 'o  
    ne', 'has', 'clue', 'going', 'crap', 'acting', 'costumes', 'can', 'get', 'acr  
    oss', 'save', 'yourself', 'an', 'hour', 'bit', 'your', 'got', 'slightly', 'an  
    sword', 'happened', 'historical', 'dream', 'based', 'based', 'historical', 'fantasy', 'the
```

```
In [ ]: train_arrays = numpy.zeros((25000, 100))

train_labels = numpy.zeros(25000)

for i in range(20000):

    train_arrays[i] = model
    train_arrays[12500 + i] = model[prefix_train_neg]
    train_labels[i] = 1
    train_labels[12500 + i] = 0
```

In []:

In []:

```
In [ ]: #model.get_latest_training_loss()
```

```
In [98]: model.wv.similarity('good','bad')
```

```
Out[98]: 0.5569832
```

```
In [99]: words = list(model.wv.vocab)
```

```
In [100]: len(words)
```

```
Out[100]: 9546
```

```
In [101]: vectors = np.array(model.wv.vectors)
```

```
In [102]: vectors
```

```
Out[102]: array([[ 0.17890231,  0.33068714, -0.8207312 , ..., -0.01275392,
                   0.3270439 ,  0.5880559 ],
                  [ 0.04594428,  0.42859146, -0.08810625, ...,  0.77135634,
                   1.0101948 ,  0.20413461],
                  [-0.5275586 ,  0.69348824,  0.42044514, ...,  0.24094042,
                   1.200559 ,  0.4792053 ],
                  ...,
                  [-1.7779503 , -1.2304244 ,  1.0064094 , ...,  0.7981529 ,
                   0.8590613 ,  1.5844318 ],
                  [ 1.0759941 ,  0.2052762 ,  1.2978704 , ...,  0.10349308,
                   1.6053356 , -0.07444835],
                  [-0.33308494,  0.4798437 ,  0.6466291 , ..., -0.5489158 ,
                   0.46562156,  1.5704718 ]], dtype=float32)
```

```
In [103]: vectors.shape
```

```
Out[103]: (9546, 100)
```

```
In [105]: model.wv.syn0
```

c:\program files\python37\lib\site-packages\ipykernel_launcher.py:1: DeprecationWarning: Call to deprecated `syn0` (Attribute will be removed in 4.0.0, use self.vectors instead).

"""Entry point for launching an IPython kernel.

```
Out[105]: array([[ 0.17890231,  0.33068714, -0.8207312 , ..., -0.01275392,
                   0.3270439 ,  0.5880559 ],
                  [ 0.04594428,  0.42859146, -0.08810625, ...,  0.77135634,
                   1.0101948 ,  0.20413461],
                  [-0.5275586 ,  0.69348824,  0.42044514, ...,  0.24094042,
                   1.200559 ,  0.4792053 ],
                  ...,
                  [-1.7779503 , -1.2304244 ,  1.0064094 , ...,  0.7981529 ,
                   0.8590613 ,  1.5844318 ],
                  [ 1.0759941 ,  0.2052762 ,  1.2978704 , ...,  0.10349308,
                   1.6053356 , -0.07444835],
                  [-0.33308494,  0.4798437 ,  0.6466291 , ..., -0.5489158 ,
                   0.46562156,  1.5704718 ]], dtype=float32)
```

```
In [106]: model.wv.syn0.shape
```

```
c:\program files\python37\lib\site-packages\ipykernel_launcher.py:1: Deprecatio
nWarning: Call to deprecated `syn0` (Attribute will be removed in 4.0.0, use se
lf.vectors instead).
    """Entry point for launching an IPython kernel.
```

```
Out[106]: (9546, 100)
```

```
In [104]: training_targets
```

```
Out[104]: array([1, 0, 0, ..., 0, 1, 0], dtype=int64)
```

```
In [ ]: training_targets.shape
```

```
In [ ]: model.save('./imdb.d2v')
```

```
In [ ]: model = Doc2Vec.load('./imdb.d2v')
```

```
In [ ]: model['']
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]: clf = LogisticRegression().fit(model.wv.syn0, training_targets[:max_dataset_size])
```

```
In [ ]: predict = clf.predict(model.wv.syn0[:100, :])
# Calculating the score of the predictions
score = clf.score(model.wv.syn0, training_targets[:max_dataset_size])
print("\nPrediction word2vec : \n", predict)
print("Score word2vec : \n", score)
```

```
In [ ]:
```

```
In [ ]: print(model['horrible'])
```

```
In [ ]: w1 = "ok"
model.wv.most_similar (positive=w1)
```

```
In [ ]: model.save('./imdb.d2v')
```

```
In [ ]: new_model = gensim.models.Word2Vec.load('./imdb.d2v')
```

- <https://machinelearningmastery.com/develop-word-embeddings-python-gensim/>
(<https://machinelearningmastery.com/develop-word-embeddings-python-gensim/>)

- <https://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.XoEUvohKiUI>
(<https://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.XoEUvohKiUI>)
- <https://stackoverflow.com/questions/49643974/how-to-do-text-classification-using-word2vec>
(<https://stackoverflow.com/questions/49643974/how-to-do-text-classification-using-word2vec>)
- https://radimrehurek.com/gensim/auto_examples/tutorials/run_doc2vec_lee.html
(https://radimrehurek.com/gensim/auto_examples/tutorials/run_doc2vec_lee.html)
- https://github.com/ibrahimsharaf/doc2vec/blob/master/models/doc2vec_model.py
(https://github.com/ibrahimsharaf/doc2vec/blob/master/models/doc2vec_model.py)